



Floating Point

- Representation for non-integral numbers
 - Including very small and very large numbers
- Like scientific notation
 - -2.34×10^{56} ← normalized
 - $+0.002 \times 10^{-4}$ ← not normalized
 - $+987.02 \times 10^9$ ← not normalized
- In binary
 - $\pm 1.xxxxxxx_2 \times 2^{yyy}$
- Types float and double in C

§3.5 Floating Point

Chapter 3 — Arithmetic for Computers — 2

Floating Point Standard

- Defined by IEEE standard 754-1985
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Two representations
 - Single precision (32-bit)
 - Double precision (64-bit)

IEEE Floating-Point Format

single: 8 bits
double: 11 bits

single: 23 bits
double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 ⇒ non-negative, 1 ⇒ negative)
- Normalized significand: $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the "1." restored
- Exponent: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1023

Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
⇒ actual exponent = 1 - 127 = -126
 - Fraction: 000...00 ⇒ significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
⇒ actual exponent = 254 - 127 = +127
 - Fraction: 111...11 ⇒ significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Chapter 3 — Arithmetic for Computers — 5

Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 00000000001
⇒ actual exponent = 1 - 1023 = -1022
 - Fraction: 000...00 ⇒ significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 11111111110
⇒ actual exponent = 2046 - 1023 = +1023
 - Fraction: 111...11 ⇒ significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Chapter 3 — Arithmetic for Computers — 6

Floating-Point Precision

- Relative precision
 - all fraction bits are significant
 - Single: approx 2^{-23}
 - Equivalent to $23 \times \log_{10} 2 \approx 23 \times 0.3$
 ≈ 6 decimal digits of precision
 - Double: approx 2^{-52}
 - Equivalent to $52 \times \log_{10} 2 \approx 52 \times 0.3$
 ≈ 16 decimal digits of precision

Chapter 3 — Arithmetic for Computers — 7

Binary Refresher

- When we look at a number like 10110_2 , we're seeing it as:
$$1(2^4) + 0(2^3) + 1(2^2) + 1(2^1) + 0(2^0)$$
$$= 16 + 4 + 2 = 22_{10}$$

Chapter 1 — Computer Abstractions and Technology — 8

Binary Decimal Points

- In decimal, 12.63 is the same as
$$1(10^1) + 2(10^0) + 6(10^{-1}) + 3(10^{-2})$$
- In binary, 101.01₂ is the same as
$$1(2^2) + 0(2^1) + 1(2^0) + 0(2^{-1}) + 1(2^{-2})$$
$$= 4 + 1 + 0.25 = 5.25$$

Useful Exponent Identities

- $a^b * a^c = a^{b+c}$
 - Why memorize more than 2¹⁰ when we can just break them down?
$$2^{35} = 2^5 * 2^{30}$$
$$= 2^5 * 2^{10} * 2^{10} * 2^{10} = 32 \text{ MB}$$
- $a^{-b} = \frac{1}{a^b}$
 - When we have decimal terms, instead of seeing 2⁻¹, 2⁻², etc. use $\frac{1}{2^1}, \frac{1}{2^2}$, etc.

More Binary

- Dividing (shifting right) by 2_{10} is the same as moving the decimal point one place to the left.
 - Same reasoning as when we divide by 10 in base 10
- Multiplying (shifting left) works the same way

Floating-Point Example

- Represent -0.75
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000...00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 0111111110_2$
- Single: $1011111101000...00$
- Double: $101111111101000...00$

Floating-Point Example

- What number is represented by the single-precision float
11000000101000...00
 - $S = 1$
 - Fraction = 01000...00₂
 - Exponent = 10000001₂ = 129
- $x = (-1)^1 \times (1 + 0.01_2) \times 2^{(129 - 127)}$
= $(-1) \times 1.25 \times 2^2$
= -5.0

IEEE 754-1985 Specials

- We reserve all 0s and all 1s in the exponent. This is why:
- 011111110000...00 = $+\infty$
- 111111110000...00 = $-\infty$
- X1111111[non-zero] = NaN
 - e.g., square root of a negative number
- X000000000000...00 = 0
 - ...there's actually a positive zero and a negative zero

Floating-Point Addition

- Consider a 4-digit decimal example
 - $9.999 \times 10^1 + 1.610 \times 10^{-1}$
- 1. Align decimal points
 - Shift number with smaller exponent
 - $9.999 \times 10^1 + 0.016 \times 10^1$
- 2. Add significands
 - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
- 3. Normalize result & check for over/underflow
 - 1.0015×10^2
- 4. Round and renormalize if necessary
 - 1.002×10^2

Chapter 3 — Arithmetic for Computers — 15

Floating-Point Addition

- Now consider a 4-digit binary example
 - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$ ($0.5 + -0.4375$)
- 1. Align binary points
 - Shift number with smaller exponent
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
 - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
 - $1.000_2 \times 2^{-4}$, with no over/underflow
- 4. Round and renormalize if necessary
 - $1.000_2 \times 2^{-4}$ (no change) = 0.0625

Chapter 3 — Arithmetic for Computers — 16

FP Instructions in MIPS

- FP hardware is coprocessor 1
 - Adjunct processor that extends the ISA
- Separate FP registers
 - 32 single-precision: \$f0, \$f1, ... \$f31
 - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
 - Release 2 of MIPS ISA supports 32 × 64-bit FP reg's
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - lwc1, ldc1, swc1, sdc1
 - e.g., ldc1 \$f8, 32(\$sp)

Chapter 3 — Arithmetic for Computers — 17

FP Instructions in MIPS

- Single-precision arithmetic
 - add.s, sub.s, mul.s, div.s
 - e.g., add.s \$f0, \$f1, \$f6
- Double-precision arithmetic
 - add.d, sub.d, mul.d, div.d
 - e.g., mul.d \$f4, \$f4, \$f6
- Single- and double-precision comparison
 - c.xx.s, c.xx.d (xx is eq, lt, le, ...)
 - Sets or clears FP condition-code bit
 - e.g. c.lt.s \$f3, \$f4
- Branch on FP condition code true or false
 - bc1t, bc1f
 - e.g., bc1t TargetLabel

Chapter 3 — Arithmetic for Computers — 18

FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

 - fahr in \$f12, result in \$f0, literals in global memory space
- Compiled MIPS code:

```
f2c: lwc1    $f16, const5($gp)  
     lwc2    $f18, const9($gp)  
     div.s   $f16, $f16, $f18  
     lwc1    $f18, const32($gp)  
     sub.s   $f18, $f12, $f18  
     mul.s   $f0,  $f16, $f18  
     jr      $ra
```

FP Example: Array Multiplication

- $X = X + Y \times Z$
 - All 32 × 32 matrices, 64-bit double-precision elements
- C code:

```
void mm (double x[][],  
         double y[][], double z[][]) {  
    int i, j, k;  
    for (i = 0; i! = 32; i = i + 1)  
        for (j = 0; j! = 32; j = j + 1)  
            for (k = 0; k! = 32; k = k + 1)  
                x[i][j] = x[i][j]  
                    + y[i][k] * z[k][j];  
}
```

 - Addresses of x, y, z in \$a0, \$a1, \$a2, and i, j, k in \$s0, \$s1, \$s2

FP Example: Array Multiplication

■ MIPS code:

	li	\$t1, 32	# \$t1 = 32 (row size/loop end)
	li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0	# j = 0; restart 2nd for loop
L2:	li	\$s2, 0	# k = 0; restart 3rd for loop
	sll	\$t2, \$s0, 5	# \$t2 = i * 32 (size of row of x)
	addu	\$t2, \$t2, \$s1	# \$t2 = i * size(row) + j
	sll	\$t2, \$t2, 3	# \$t2 = byte offset of [i][j]
	addu	\$t2, \$a0, \$t2	# \$t2 = byte address of x[i][j]
	l.d	\$f4, 0(\$t2)	# \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5	# \$t0 = k * 32 (size of row of z)
	addu	\$t0, \$t0, \$s1	# \$t0 = k * size(row) + j
	sll	\$t0, \$t0, 3	# \$t0 = byte offset of [k][j]
	addu	\$t0, \$a2, \$t0	# \$t0 = byte address of z[k][j]
	l.d	\$f16, 0(\$t0)	# \$f16 = 8 bytes of z[k][j]

...

FP Example: Array Multiplication

	sll	\$t0, \$s0, 5	# \$t0 = i*32 (size of row of y)
	addu	\$t0, \$t0, \$s2	# \$t0 = i*size(row) + k
	sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
	addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
	l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
	mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
	add.d	\$f4, \$f4, \$f16	# f4=x[i][j] + y[i][k]*z[k][j]
	addiu	\$s2, \$s2, 1	# \$k k + 1
	bne	\$s2, \$t1, L3	# if (k != 32) go to L3
	s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
	addiu	\$s1, \$s1, 1	# \$j = j + 1
	bne	\$s1, \$t1, L2	# if (j != 32) go to L2
	addiu	\$s0, \$s0, 1	# \$i = i + 1
	bne	\$s0, \$t1, L1	# if (i != 32) go to L1

Accurate Arithmetic

- IEEE Std 754 specifies additional rounding control
 - Extra bits of precision (guard, round, sticky)
 - Choice of rounding modes
 - Allows programmer to fine-tune numerical behavior of a computation
- Not all FP units implement all options
 - Most programming languages and FP libraries just use defaults
- Trade-off between hardware complexity, performance, and market requirements

Associativity

- Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38		-1.50E+38
y	1.50E+38	0.00E+00	
z	1.0	1.0	1.50E+38
		1.00E+00	0.00E+00

- Need to validate parallel programs under varying degrees of parallelism